

recipe_recommender_hybrid_model-Copy1

November 30, 2023

```
[51]: import numpy as np
import pandas as pd
import pickle

# Load the datasets
ingr_map = pd.read_pickle("ingr_map.pkl")
raw_recipes = pd.read_csv("RAW_recipes.csv")
raw_interactions = pd.read_csv("RAW_interactions.csv")
pp_users = pd.read_csv("PP_users.csv")
pp_recipes = pd.read_csv("PP_recipes.csv")
interactions_validation = pd.read_csv("interactions_validation.csv")
interactions_train = pd.read_csv("interactions_train.csv")
interactions_test = pd.read_csv("interactions_test.csv")

# Show some basic information about each dataset
datasets = {
    'ingr_map': ingr_map,
    'raw_recipes': raw_recipes,
    'raw_interactions': raw_interactions,
    'pp_users': pp_users,
    'pp_recipes': pp_recipes,
    'interactions_validation': interactions_validation,
    'interactions_train': interactions_train,
    'interactions_test': interactions_test
}

info_dict = {}

for name, dataset in datasets.items():
    info_dict[name] = {
        'Number of Rows': dataset.shape[0],
        'Number of Columns': dataset.shape[1],
        'Columns': ', '.join(dataset.columns)
    }

info_df = pd.DataFrame(info_dict).T
```

```
info_df
```

```
[51]:
```

	Number of Rows	Number of Columns	\
ingr_map	11659	7	
raw_recipes	231637	12	
raw_interactions	1132367	5	
pp_users	25076	6	
pp_recipes	178265	8	
interactions_validation	7023	6	
interactions_train	698901	6	
interactions_test	12455	6	

	Columns
ingr_map	raw_ingr, raw_words, processed, len_proc, repl...
raw_recipes	name, id, minutes, contributor_id, submitted, ...
raw_interactions	user_id, recipe_id, date, rating, review
pp_users	u, techniques, items, n_items, ratings, n_ratings
pp_recipes	id, i, name_tokens, ingredient_tokens, steps_t...
interactions_validation	user_id, recipe_id, date, rating, u, i
interactions_train	user_id, recipe_id, date, rating, u, i
interactions_test	user_id, recipe_id, date, rating, u, i

```
[22]: raw_recipes_df = pd.read_csv('RAW_recipes.csv')
raw_interactions_df = pd.read_csv("RAW_interactions.csv")
```

1 SVD collaborative filtering

```
[23]: from surprise import Dataset, Reader
from surprise import SVD
from surprise.model_selection import cross_validate
from surprise.model_selection import train_test_split
from surprise import accuracy
```

```
[24]: data = raw_interactions_df[['user_id', 'recipe_id', 'rating']]
```

```
[5]: reader = Reader(rating_scale=(1, 5))
data = Dataset.load_from_df(data, reader)
model = SVD()

cross_validation_results = cross_validate(model, data, measures=['RMSE', 'MAE'], cv=5, verbose=True)
print(cross_validation_results)
```

Evaluating RMSE, MAE of algorithm SVD on 5 split(s).

	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Mean	Std
RMSE (testset)	1.2236	1.2252	1.2139	1.2206	1.2193	1.2205	0.0039

```

MAE (testset)      0.7417  0.7416  0.7362  0.7394  0.7397  0.7397  0.0020
Fit time           12.16   12.13   12.30   12.57   11.94   12.22   0.21
Test time          1.44    1.37    1.32    1.41    1.67    1.44    0.12
{'test_rmse': array([1.22363968, 1.22522626, 1.21388836, 1.22058748,
1.21926462]), 'test_mae': array([0.74166363, 0.7416333 , 0.73620976, 0.73935663,
0.73973383]), 'fit_time': (12.156907081604004, 12.127675294876099,
12.29780387878418, 12.570320844650269, 11.936063051223755), 'test_time':
(1.437319040298462, 1.3681721687316895, 1.316875696182251, 1.4091508388519287,
1.6688730716705322)}

```

```

[6]: # Split the data into training and test set (e.g., 75% training, 25% testing)
trainset, testset = train_test_split(data, test_size=0.25)

# Train the model on the training set
model = SVD()
model.fit(trainset)

# Make predictions on the test set
predictions = model.test(testset)

# Compute and print the accuracy metrics
rmse = accuracy.rmse(predictions)
mae = accuracy.mae(predictions)

```

RMSE: 1.2214

MAE: 0.7404

2 Hyperparameter tuning for SVD collaborative filtering

```

[25]: from surprise import SVD
from surprise.model_selection import GridSearchCV
from surprise import Dataset, Reader

```

```

[7]: # Define the parameter grid
param_grid = {
    'n_factors': [50, 100, 150],
    'n_epochs': [20, 30, 40],
    'lr_all': [0.002, 0.005],
    'reg_all': [0.02, 0.1]
}

# Setup grid search
gs = GridSearchCV(SVD, param_grid, measures=['rmse'], cv=3)

# Load the dataset
reader = Reader(rating_scale=(1, 5))

```

```

data = Dataset.load_from_df(raw_interactions_df[['user_id', 'recipe_id',
↪ 'rating']], reader)

# Run grid search
gs.fit(data)

print("Best RMSE score: ", gs.best_score['rmse'])
print("Best parameters: ", gs.best_params['rmse'])

```

Best RMSE score: 1.2161208333623452
Best parameters: {'n_factors': 50, 'n_epochs': 20, 'lr_all': 0.005, 'reg_all': 0.1}

3 Using the Best Parameters to Train the SVD collaborative filtering Model

```

[26]: # Setup the SVD model with the best parameters
optimized_SVD = SVD(n_factors=50, n_epochs=20, lr_all=0.005, reg_all=0.1)

reader = Reader(rating_scale=(1, 5))
data = Dataset.load_from_df(raw_interactions_df[['user_id', 'recipe_id',
↪ 'rating']], reader)

# Split your dataset into train and test sets
trainset, testset = train_test_split(data, test_size=0.25)

# Train the model on the trainset
optimized_SVD.fit(trainset)

# Make predictions on the testset
predictions = optimized_SVD.test(testset)

# Calculate and print the RMSE on the test set
rmse = accuracy.rmse(predictions)
mae = accuracy.mae(predictions)

```

RMSE: 1.2155
MAE: 0.7407

```

[28]: cross_validation_optimized_SVD = cross_validate(optimized_SVD, data,
↪ measures=['RMSE', 'MAE'], cv=5, verbose=True)
print(cross_validation_optimized_SVD)

```

Evaluating RMSE, MAE of algorithm SVD on 5 split(s).

	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Mean	Std
RMSE (testset)	1.2099	1.2150	1.2124	1.2163	1.2117	1.2130	0.0023

```
MAE (testset)      0.7381  0.7407  0.7399  0.7421  0.7398  0.7401  0.0013
Fit time          9.98    9.80    9.83    10.45   10.56   10.12   0.32
Test time         2.78    2.39    2.51    2.52    2.54    2.55    0.13
{'test_rmse': array([1.20986348, 1.21495517, 1.21239871, 1.21627694,
1.21166121]), 'test_mae': array([0.73814761, 0.74072785, 0.73988309, 0.74213283,
0.73976056]), 'fit_time': (9.978832960128784, 9.798351049423218,
9.83065915107727, 10.449662208557129, 10.555390119552612), 'test_time':
(2.7790751457214355, 2.3928768634796143, 2.5148427486419678, 2.5208659172058105,
2.5412588119506836)}
```

4 Co-clustering Collaborative Filtering Model

```
[6]: from surprise import CoClustering

reader = Reader(rating_scale=(1, 5))
data = Dataset.load_from_df(raw_interactions_df[['user_id', 'recipe_id',
↪ 'rating']], reader)

trainset, testset = train_test_split(data, test_size=0.25)
co_clustering_model = CoClustering()
co_clustering_model.fit(trainset)
```

```
[6]: <surprise.prediction_algorithms.co_clustering.CoClustering at 0x7fac2163fa00>
```

```
[7]: predictions = co_clustering_model.test(testset)
rmse = accuracy.rmse(predictions)
mae = accuracy.mae(predictions)
```

```
RMSE: 1.3097
```

```
MAE: 0.7562
```

```
[8]: param_grid = {'n_cltr_u': [3, 5, 7], 'n_cltr_i': [3, 5, 7], 'n_epochs': [20,
↪ 30, 40]}
gs = GridSearchCV(CoClustering, param_grid, measures=['rmse', 'mae'], cv=3)
gs.fit(data)

print("Best RMSE score: ", gs.best_score['rmse'])
print("Best parameters: ", gs.best_params['rmse'])
```

```
Best RMSE score: 1.3103096749322833
```

```
Best parameters: {'n_cltr_u': 3, 'n_cltr_i': 5, 'n_epochs': 20}
```

```
[9]: # Initialize the Co-clustering model with the best parameters
model = CoClustering(n_cltr_u=3, n_cltr_i=5, n_epochs=20)

# Train the model on the entire dataset
trainset, testset = train_test_split(data, test_size=0.25)
```

```
model.fit(trainset)
```

```
[9]: <surprise.prediction_algorithms.co_clustering.CoClustering at 0x7fac2165e130>
```

```
[10]: predictions = model.test(testset)
      rmse = accuracy.rmse(predictions)
      mae = accuracy.mae(predictions)
```

RMSE: 1.3041

MAE: 0.7479

5 Content-based Recommendation Model

```
[13]: # Get Tags for Content-based features
      import ast

      recipe_data = pd.read_csv("RAW_recipes.csv")

      # Parsing the tags from string representation of list to actual list
      recipe_data['tags'] = recipe_data['tags'].apply(ast.literal_eval)

      # Exploring the unique tags and their frequencies
      all_tags = [tag for sublist in recipe_data['tags'] for tag in sublist]
      unique_tags = set(all_tags)
      tag_frequency = pd.Series(all_tags).value_counts()

      num_unique_tags = len(unique_tags)
      selected_indices = [4,5,8, 11] + list(range(13, 61))
      selected_tags = tag_frequency.iloc[selected_indices]

      num_unique_tags, selected_tags
```

```
[13]: (552,
      dietary          165091
      easy             126062
      low-in-something  85776
      60-minutes-or-less 69990
      meat             56042
      30-minutes-or-less 55077
      vegetables       53814
      taste-mood       52143
      4-hours-or-less  49497
      north-american   48479
      3-steps-or-less  44933
      15-minutes-or-less 43934
      low-sodium       43349)
```

desserts	43203
low-carb	42189
healthy	40340
dinner-party	37561
low-cholesterol	36743
low-calorie	36429
vegetarian	35651
beginner-cook	35561
5-ingredients-or-less	35466
holiday-event	34920
inexpensive	32619
low-protein	32522
low-saturated-fat	31378
fruit	31324
oven	31180
american	31179
eggs-dairy	30142
pasta-rice-and-grains	27084
kid-friendly	27074
side-dishes	26902
healthy-2	26619
comfort-food	26136
european	24912
presentation	24470
poultry	24160
lunch	23800
for-1-or-2	23084
low-fat	22170
stove-top	22095
seasonal	21933
weeknight	20948
chicken	20381
appetizers	20379
brunch	18927
to-go	18524
for-large-groups	17391
beef	17074
one-dish-meal	16807
cheese	15147

Name: count, dtype: int64)

```
[14]: # One-hot encoding tags to speed up computation
top_tags = selected_tags.index.tolist()

# Initializing columns for top tags with default value 0
for tag in top_tags:
    recipe_data[f'tag_{tag}'] = 0
```

```

# Setting the value to 1 if the recipe contains the tag
for index, row in recipe_data.iterrows():
    for tag in top_tags:
        if tag in row['tags']:
            recipe_data.at[index, f'tag_{tag}'] = 1

```

```

[15]: # Merging the user ratings data (interactions_data) with the one-hot encoded
      ↪ tags from recipe_data
      # The merging key will be 'recipe_id'

interactions_data = pd.read_csv("RAW_interactions.csv")

# Selecting relevant columns from recipe_data (recipe_id and one-hot encoded
      ↪ tags)
recipe_tags_data = recipe_data[['id'] + [col for col in recipe_data.columns if
      ↪ col.startswith('tag_')]]

# Renaming 'id' column to 'recipe_id' for consistency
recipe_tags_data.rename(columns={'id': 'recipe_id'}, inplace=True)

# Merging the datasets
merged_data = interactions_data.merge(recipe_tags_data, how='left',
      ↪ on='recipe_id')

```

/var/folders/bg/hdj9jw_j33g1vg7x_rmvr9lc0000gn/T/ipykernel_10539/1372912005.py:1

0: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
recipe_tags_data.rename(columns={'id': 'recipe_id'}, inplace=True)
```

```

[16]: # Filling missing values in tag columns with zeros
merged_data.fillna({col: 0 for col in merged_data.columns if col.
      ↪ startswith('tag_')}, inplace=True)

# Checking the first few rows of the updated merged dataset
merged_data.head()

```

```

[16]:   user_id  recipe_id      date  rating  \
0    38094    40893  2003-02-17        4
1   1293707    40893  2011-12-21        5
2     8937    44394  2002-12-01        4
3   126440    85009  2010-02-27        5
4    57222    85009  2011-10-01        5

```


		review	tag_dietary	tag_easy	\
0	Great with a salad. Cooked on top of stove for...		1	1	
1	So simple, so delicious! Great for chilly fall...		1	1	
2	This worked very well and is EASY. I used not...		1	1	
3	I made the Mexican topping and took it to bunk...		0	1	
4	Made the cheddar bacon topping, adding a sprin...		0	1	

	tag_low-in-something	tag_60-minutes-or-less	tag_meat	...	tag_seasonal	\
0	0	0	0	...	0	
1	0	0	0	...	0	
2	0	0	0	...	1	
3	0	0	0	...	0	
4	0	0	0	...	0	

	tag_weeknight	tag_chicken	tag_appetizers	tag_brunch	tag_to-go	\
0	1	0	0	0	0	
1	1	0	0	0	0	
2	0	0	0	0	1	
3	0	0	0	0	0	
4	0	0	0	0	0	

	tag_for-large-groups	tag_beef	tag_one-dish-meal	tag_cheese
0	0	0	0	0
1	0	0	0	0
2	1	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 57 columns]

```
[35]: merged_data2= merged_data
```

6 Hyperparameter Tuning for Content-based Model

```
[17]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor

# Prepare the dataset for content-based model
X = merged_data.drop(columns=['user_id', 'recipe_id', 'rating', 'review',
↪ 'date'])
y = merged_data['rating']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Train RandomForest model
```

```

content_model = RandomForestRegressor()
content_model.fit(X_train, y_train)

# Make predictions on the test set
predictions_content = content_model.predict(X_test)

```

```

[18]: from sklearn.metrics import mean_squared_error, mean_absolute_error

mse = mean_squared_error(y_test, predictions_content)
print("Mean Squared Error (MSE):", mse)

rmse = mean_squared_error(y_test, predictions_content, squared=False)
print("Root Mean Squared Error (RMSE):", rmse)

mae = mean_absolute_error(y_test, predictions_content)
print("Mean Absolute Error (MAE):", mae)

```

Mean Squared Error (MSE): 1.6841104017151067
Root Mean Squared Error (RMSE): 1.2977327928796076
Mean Absolute Error (MAE): 0.8311654118787669

```

[ ]: #Optimizing hyperparameters
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestRegressor

# Define the parameter grid
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [10, 20, 30],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}

# Create a base model
rf = RandomForestRegressor()

# Instantiate the grid search model
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=3,
    ↪n_jobs=-1, verbose=2)

# Fit the grid search to the data
grid_search.fit(X_train, y_train)

# Best parameters
best_params = grid_search.best_params_
print("Best parameters:", best_params)

```

Best parameters: {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_split': 5,

```
‘n_estimators’: 300}
```

```
[40]: # Create a new model with the best parameters
optimized_rf = RandomForestRegressor(max_depth=10, min_samples_leaf=2,
    ↪min_samples_split=5, n_estimators=300)

# Retrain the model on the entire training set
optimized_rf.fit(X_train, y_train)

# Predict on the test set
y_pred = optimized_rf.predict(X_test)

# Calculate evaluation metrics
mse = mean_squared_error(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred, squared=False)
mae = mean_absolute_error(y_test, y_pred)

print("MSE:", mse)
print("RMSE:", rmse)
print("MAE:", mae)
```

```
MSE: 1.6006633320865211
RMSE: 1.2651732419263857
MAE: 0.8456339830113258
```

```
[29]: # Running cross-validation

from sklearn.model_selection import cross_val_score

optimized_rf = RandomForestRegressor(max_depth=10, min_samples_leaf=2,
    ↪min_samples_split=5, n_estimators=300)
num_folds = 5

mse_scores = cross_val_score(optimized_rf, X, y,
    ↪scoring='neg_mean_squared_error', cv=num_folds)
mse_scores = -mse_scores
print("Mean MSE:", mse_scores.mean())

rmse_scores = cross_val_score(optimized_rf, X, y,
    ↪scoring='neg_root_mean_squared_error', cv=num_folds)
rmse_scores = -rmse_scores
print("Mean RMSE:", rmse_scores.mean())

mae_scores = cross_val_score(optimized_rf, X, y,
    ↪scoring='neg_mean_absolute_error', cv=num_folds)
mae_scores = -mae_scores
print("Mean MAE:", mae_scores.mean())
```

Mean MSE: 1.5944072085619527
Mean RMSE: 1.2627100066704786
Mean MAE: 0.8473926819010259

7 Creating a Hybrid Model with collaborative filtering and content-based model

```
[31]: # Train the models individually
# RandomForest
optimized_rf = RandomForestRegressor(max_depth=10, min_samples_leaf=2,
    ↪min_samples_split=5, n_estimators=300)

X = merged_data.drop(columns=['user_id', 'recipe_id', 'rating', 'review',
    ↪'date'])
y = merged_data['rating']
#X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

optimized_rf.fit(X_train, y_train)

# SVD
optimized_SVD = SVD(n_factors=50, n_epochs=20, lr_all=0.005, reg_all=0.1)

reader = Reader(rating_scale=(1, 5))
data = Dataset.load_from_df(raw_interactions_df[['user_id', 'recipe_id',
    ↪'rating']], reader)
trainset = data.build_full_trainset()
optimized_SVD.fit(trainset)
```

```
[31]: <surprise.prediction_algorithms.matrix_factorization.SVD at 0x7fab35d8f280>
```

```
[54]: # Optimizing model weights
from sklearn.model_selection import train_test_split

feature_columns = [col for col in merged_data2.columns if col.
    ↪startswith('tag_')]
X = merged_data2[feature_columns]
y = merged_data2['rating']

user_ids = merged_data2['user_id']
item_ids = merged_data2['recipe_id']

X_train, X_test, y_train, y_test, user_ids_train, user_ids_test,
    ↪item_ids_train, item_ids_test = train_test_split(
    X, y, user_ids, item_ids, test_size=0.2, random_state=42)

reader = Reader(rating_scale=(1, 5))
```

```

# Create the training dataset for SVD
train_data = pd.DataFrame({
    'user_id': user_ids_train,
    'item_id': item_ids_train,
    'rating': y_train
})
train_data = Dataset.load_from_df(train_data, reader)
trainset = train_data.build_full_trainset()

def weighted_prediction(user_id, item_id, features, weight_rf, weight_svd):
    rf_prediction = optimized_rf.predict([features])[0]
    svd_prediction = optimized_SVD.predict(user_id, item_id).est
    return weight_rf * rf_prediction + weight_svd * svd_prediction

def compute_error(weights, user_ids, item_ids, features, actual_ratings):
    weight_rf, weight_svd = weights
    predictions = [
        weighted_prediction(user_id, item_id, feature, weight_rf, weight_svd)
        for user_id, item_id, feature in zip(user_ids, item_ids, features)
    ]
    return mean_squared_error(actual_ratings, predictions)

```

[]:

```

[ ]: from scipy.optimize import minimize

subset_size = 10000
subset_indices = np.random.choice(X_train.index, subset_size, replace=False)

X_subset = X_train.loc[subset_indices].values
y_subset = y_train.loc[subset_indices].values
user_ids_subset = user_ids_train.loc[subset_indices].values
item_ids_subset = item_ids_train.loc[subset_indices].values

# Initial guesses for weights
initial_weights = [0.5, 0.5]

# The bounds ensure that weights are between 0 and 1
bounds = [(0, 1), (0, 1)]

# Perform the optimization with a tolerance value
result = minimize(
    compute_error,
    initial_weights,
    args=(user_ids_subset, item_ids_subset, X_subset, y_subset),
    bounds=bounds,

```

```

        method='SLSQP', # Sequential Least Squares Programming
        tol=1e-3 # Adjust the tolerance for faster convergence
    )

    optimized_weights = result.x
    print("Optimized weights:", optimized_weights)

```

```
[57]: print("Optimized weights:", optimized_weights)
```

Optimized weights: [0. 1.]

```

[ ]: # Testing the model with optimized weights

X_test_rf = merged_data2.drop(columns=['user_id', 'recipe_id', 'date',
    ↪ 'rating', 'review']) # Prepare the data for the RandomForest model

test_data_svd = merged_data2[['user_id', 'recipe_id']] # Prepare the data for
    ↪ the SVD model

y_true = merged_data2['rating'] # Actual ratings for evaluation

weight_rf = 0
weight_svd = 1

hybrid_predictions = []
for index, row in merged_data2.iterrows():
    # Get features for RandomForest
    features_rf = row.drop(['user_id', 'recipe_id', 'date', 'rating',
    ↪ 'review']).values.reshape(1, -1)

    # Make predictions using both models
    rf_pred = optimized_rf.predict(features_rf)[0]
    svd_pred = optimized_SVD.predict(str(row['user_id']),
    ↪ str(row['recipe_id'])).est

    # Combine predictions using a weighted average
    hybrid_pred = (rf_pred * weight_rf) + (svd_pred * weight_svd)
    hybrid_predictions.append(hybrid_pred)

# Evaluate the hybrid model
from sklearn.metrics import mean_squared_error, mean_absolute_error

mse = mean_squared_error(y_true, hybrid_predictions)
rmse = mean_squared_error(y_true, hybrid_predictions, squared=False)

```

```
mae = mean_absolute_error(y_true, hybrid_predictions)
```

```
print(f'MSE: {mse}, RMSE: {rmse}, MAE: {mae}')
```

```
[59]: print(f'MSE: {mse}, RMSE: {rmse}, MAE: {mae}')
```

```
MSE: 1.5995958057022754, RMSE: 1.2647512821508724, MAE: 0.8492393396162298
```

```
[ ]:
```